



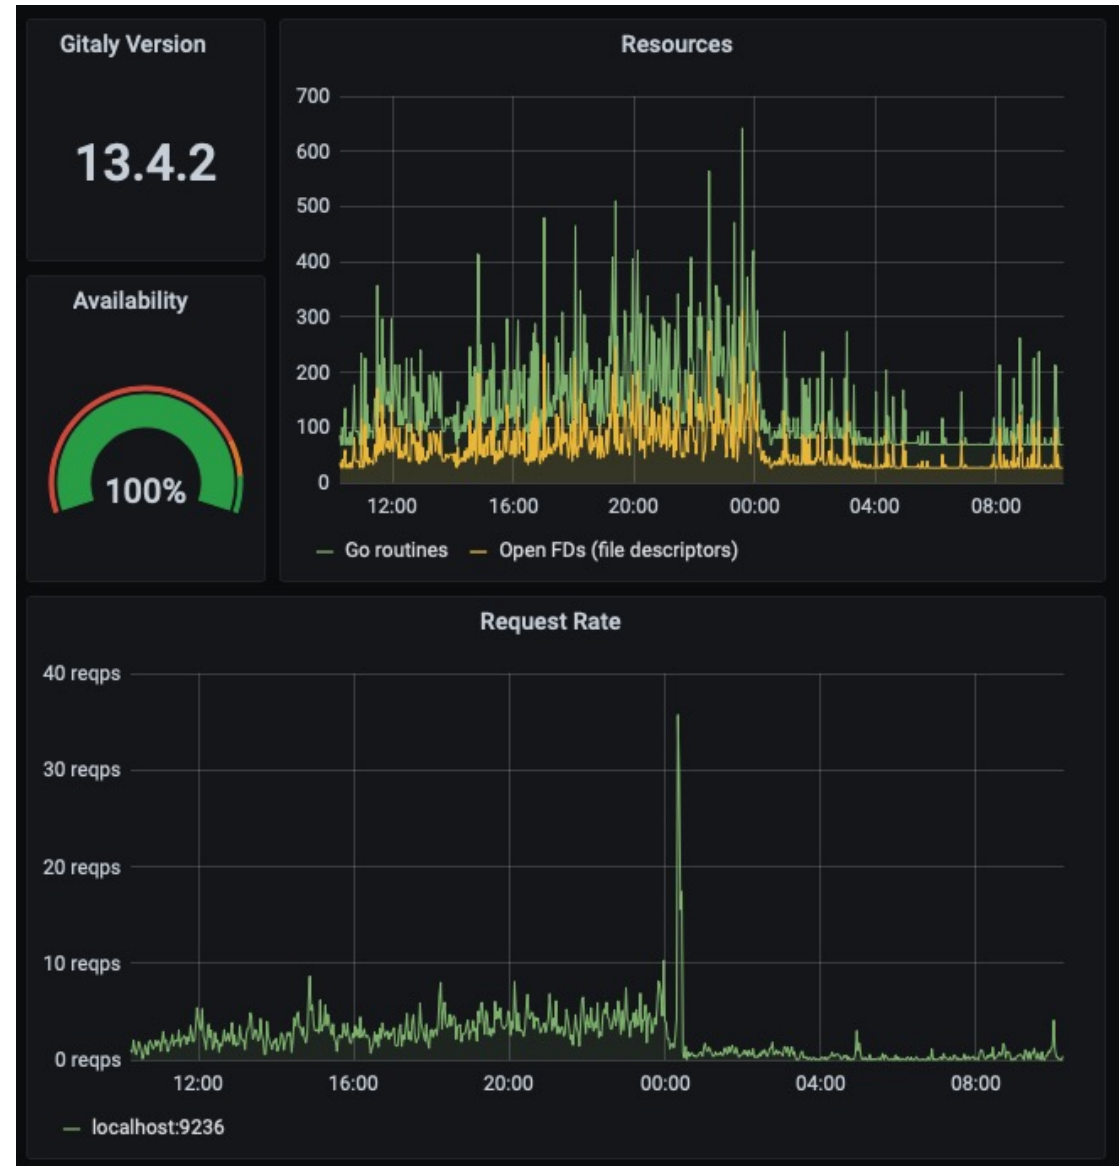
Graphs

Prof. Darrell Long

CSE 13S

Is this a graph?


- Yes, this is commonly called a graph, but it is not what we are talking about here...
- We will call these depictions of numerical quantities *plots*.



Is this a graph?

- This is an example of what we will call a *graph*.
- There is an entire branch of mathematics dedicated to *graph theory*.
- Graphs of this type are used in mathematics, computer science, physics, and even sociology.


Vertices and Edges



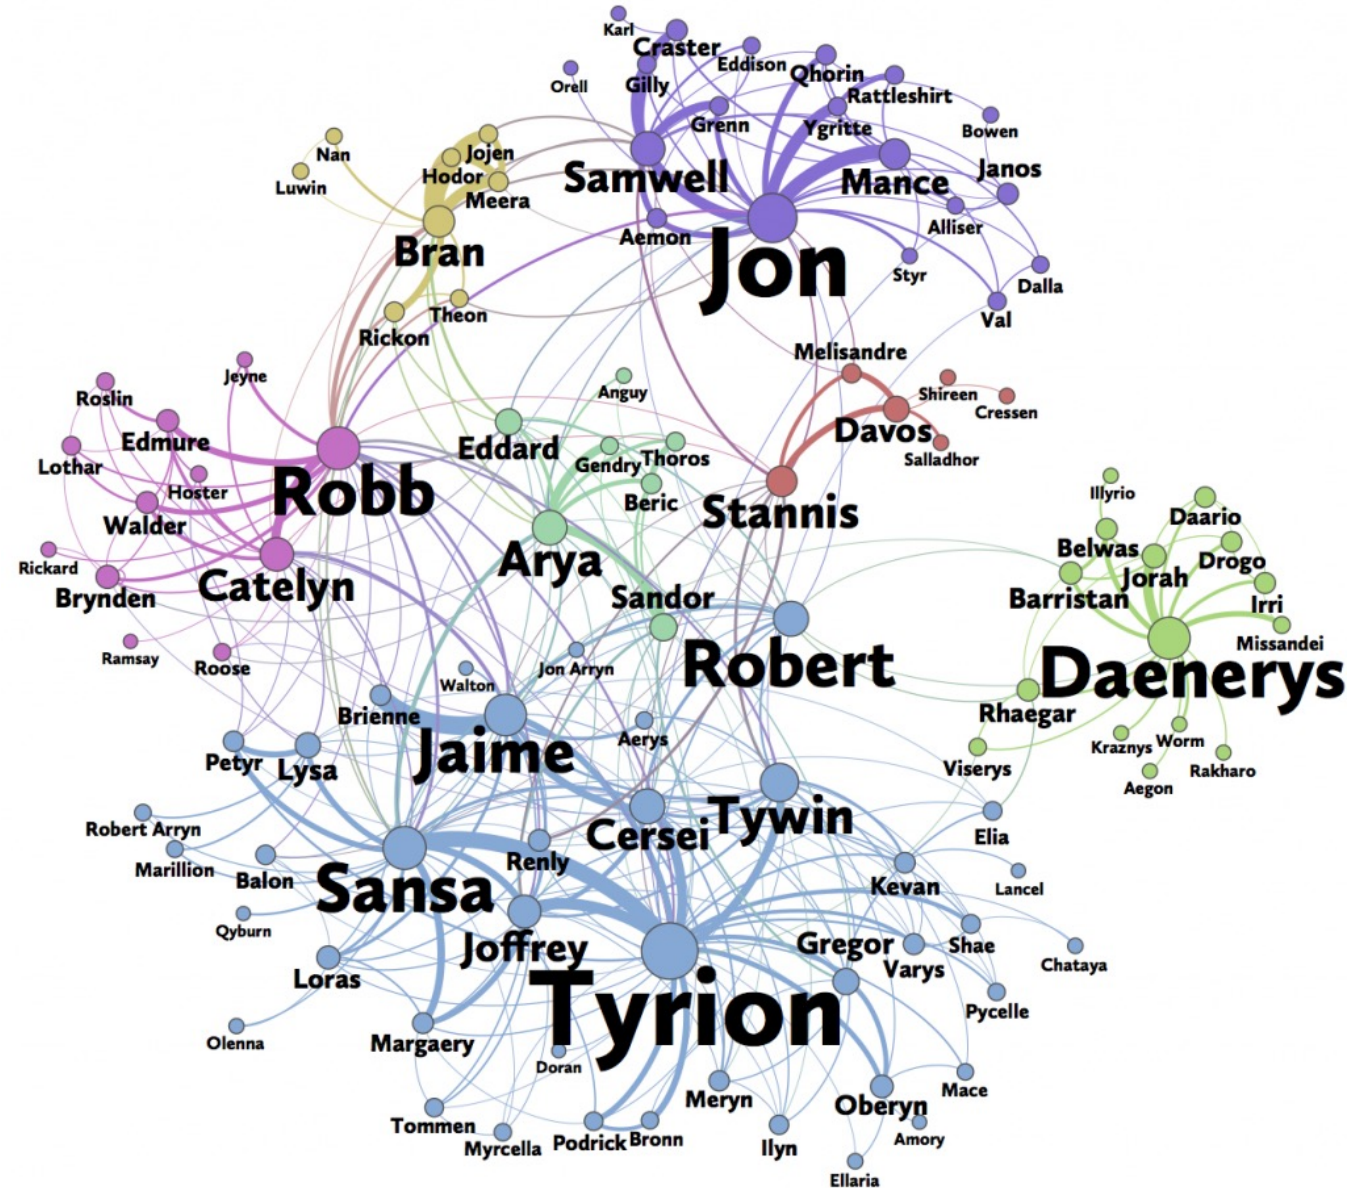
Network Routing

- The Internet (originally called ARPAnet) is a graph.
 - The nodes are computers (one reason that we call computers nodes).
 - The edges are the connections.
- Routers are **nodes** with many **edges**.

Social Networks



Network of Throne



A Formal Definition


- We formally define a graph G as follows:
 - $G = \langle V, E \rangle$
 - A graph is defined by its **vertices** and **edges**.
- **V** is the set of vertices.
 - $V = \{v_1, v_2, \dots, v_n\}$
- **E** is the set of edges.
 - Each edge is a tuple of vertices.
 - $E = \{\langle v_i, v_j \rangle, \langle v_p, v_q \rangle, \dots, \langle v_s, v_t \rangle\}$

Directed and Undirected Graphs

- Edges may have a direction, $n_1 \rightarrow n_2$, and we call that a *directed graph*.
- Edges may have no direction (or both directions), $n_1 \leftrightarrow n_2$, and we call that an *undirected graph*.
- The edges *may have weights*, which represent capacity, strength, or cost.

Representing a graph


- Adjacency Matrix
 - $n \times n$ matrix
 - Binary: edges present or absent
 - Weighted: $n \neq 0$
- Adjacency List
 - Column array for nodes
 - Linked list of edges from each node
 - May contain weights



Computed by Wolfram|Alpha

Adjacency Matrix (Used in Assignment 4)

- A non-zero entry in $M_{i,j}$ means there is an edge $n_i \rightarrow n_j$.
- A matrix that is symmetric around the diagonal represents an *undirected* graph.
- The entry can specify not only the existence of an edge, but also its weight.
- Requires $O(n^2)$ space.
 - Sparse matrix techniques can improve it.



Computed by Wolfram|Alpha

A graph in C

- Uses dynamically sized adjacency matrix.
- Adding an edge is $O(1)$.
- Checking for the existence of an edge is also $O(1)$.

```
typedef struct {
    uint32_t vertices; // Number of vertices in graph.
    uint32_t **matrix; // Adjacency matrix.
} Graph;

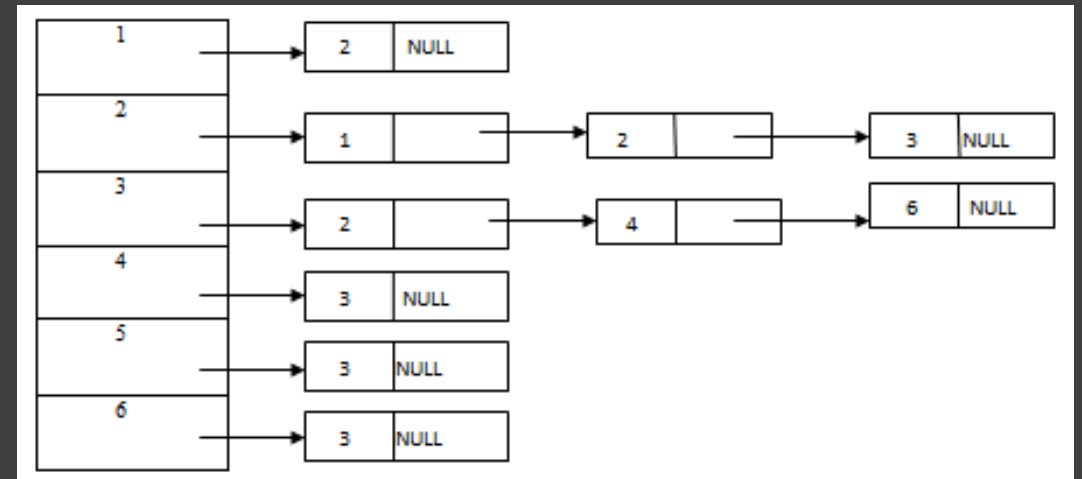
// Constructor for a graph with dynamically sized adjacency matrix.
// Performs no error checking with memory allocation.
Graph *graph_create(uint32_t vertices) {
    Graph *g = (Graph *)malloc(sizeof(Graph));
    g->vertices = vertices;
    g->matrix = (uint32_t **)calloc(vertices, sizeof(uint32_t *));
    for (uint32_t i = 0; i < vertices; i += 1) {
        g->matrix[i] = (uint32_t *)calloc(vertices, sizeof(uint32_t));
    }
    return g;
}

// Adds an edge <i, j> of weight k.
// Does not check if vertices are valid.
void graph_add_edge(Graph *g, uint32_t i, uint32_t j, uint32_t k) {
    g->matrix[i][j] = k;
}

// Checks if edge <i, j> exists.
// Does not check if vertices are valid.
bool graph_has_edge(Graph *g, uint32_t i, uint32_t j) {
    return g->matrix[i][j] > 0;
}
```


Adjacency List

- Each node is represented as an entry in a column vector.
 - Each entry is the head of a linked list.
- The list elements contain:
 - The destination node, and
 - The *weight* of the edge.
- Why would you prefer this over an adjacency matrix?
 - An adjacency matrix is $O(n^2)$ space,
 - An adjacency list will be more space efficient for sparse graphs.



Another graph in C

- Uses adjacency lists.
 - Implemented using *linked lists*.
- Adding an edge is $O(1)$.
- Checking for the existence of an edge is $O(n)$.
 - Must traverse entire list!

```
#include "ll.h" // Linked list ADT interface.

typedef struct {
    uint32_t vertices; // Number of vertices in graph.
    LinkedList **lists; // Adjacency lists.
} Graph;

// Constructor for a graph using adjacency lists.
// Performs no error checking with memory allocation.
Graph *graph_create(uint32_t vertices) {
    Graph *g = (Graph *)malloc(sizeof(Graph));
    g->vertices = vertices;
    g->lists = (LinkedList **)calloc(vertices, sizeof(LinkedList *));
    for (uint32_t i = 0; i < vertices; i += 1) {
        g->lists[i] = ll_create(); // Initialize each to be empty adjacency list.
    }
    return g;
}

// Adds an edge <i, j> of weight k.
// Does not check if vertices are valid.
void graph_add_edge(Graph *g, uint32_t i, uint32_t j, uint32_t k) {
    // Inserts a new node containing j as the vertex and k as the weight.
    ll_insert(g->lists[i], j, k);
}

// Checks if edge <i, j> exists.
// Does not check if vertices are valid.
bool graph_has_edge(Graph *g, uint32_t i, uint32_t j) {
    // Searches the list, looking for a node that contains j as the vertex.
    return ll_find(g->lists[i], j);
}
```

Basic Graph Algorithms

- We often want to find something in a graph.
- Two ways of searching a graph:

1. Breadth-first search

- Uses a queue.
- Explore the set of vertices immediately reachable.
 - Then repeat the process for each vertex in the set.
- Also known as “level-order” traversal.

2. Depth-first search

- Uses recursion or a stack.
- Search as far as possible before backing up.
- We will showcase *iterative* DFS using a stack.

Depth-First Search


- With a Tree
 - Basic recursive traversal
 - Visit all of the nodes
- With a Graph
 - Can be recursive
 - But need to mark which nodes have been visited
 - The recursive algorithm's "Base Case"

Depth-First-Search (DFS)

```
void dfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Track where we've been.
    Stack *s = stack_create(graph_vertices(g)); // Capacity of stack is number of vertices.
    stack_push(s, v);

    while (!stack_empty(s)) {
        stack_pop(s, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                stack_push(s, u);
            }
        }
    }
}
```



Stack: 0


Top of stack will be the rightmost element.

Depth-First-Search (DFS)

```
void dfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Track where we've been.
    Stack *s = stack_create(graph_vertices(g)); // Capacity of stack is number of vertices.
    stack_push(s, v);

    while (!stack_empty(s)) {
        stack_pop(s, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                stack_push(s, u);
            }
        }
    }
}
```



Stack: 1, 2, 3


Top of stack will be the rightmost element.

Depth-First-Search (DFS)

```
void dfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Track where we've been.
    Stack *s = stack_create(graph_vertices(g)); // Capacity of stack is number of vertices.
    stack_push(s, v);

    while (!stack_empty(s)) {
        stack_pop(s, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                stack_push(s, u);
            }
        }
    }
}
```



Stack: 1, 2, 7


Top of stack will be the rightmost element.

Depth-First-Search (DFS)

```
void dfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Track where we've been.
    Stack *s = stack_create(graph_vertices(g)); // Capacity of stack is number of vertices.
    stack_push(s, v);

    while (!stack_empty(s)) {
        stack_pop(s, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                stack_push(s, u);
            }
        }
    }
}
```



Stack: 1, 2


Top of stack will be the rightmost element.

Depth-First-Search (DFS)

```
void dfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Track where we've been.
    Stack *s = stack_create(graph_vertices(g)); // Capacity of stack is number of vertices.
    stack_push(s, v);

    while (!stack_empty(s)) {
        stack_pop(s, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                stack_push(s, u);
            }
        }
    }
}
```



Stack: 1, 6


Top of stack will be the rightmost element.

Depth-First-Search (DFS)

```
void dfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Track where we've been.
    Stack *s = stack_create(graph_vertices(g)); // Capacity of stack is number of vertices.
    stack_push(s, v);

    while (!stack_empty(s)) {
        stack_pop(s, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                stack_push(s, u);
            }
        }
    }
}
```



Stack: 1


Top of stack will be the rightmost element.

Depth-First-Search (DFS)

```
void dfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Track where we've been.
    Stack *s = stack_create(graph_vertices(g)); // Capacity of stack is number of vertices.
    stack_push(s, v);

    while (!stack_empty(s)) {
        stack_pop(s, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                stack_push(s, u);
            }
        }
    }
}
```



Stack: 4, 5


Top of stack will be the rightmost element.

Depth-First-Search (DFS)

```
void dfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Track where we've been.
    Stack *s = stack_create(graph_vertices(g)); // Capacity of stack is number of vertices.
    stack_push(s, v);

    while (!stack_empty(s)) {
        stack_pop(s, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                stack_push(s, u);
            }
        }
    }
}
```



Stack: 4


Top of stack will be the rightmost element.

Depth-First-Search (DFS)

```
void dfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Track where we've been.
    Stack *s = stack_create(graph_vertices(g)); // Capacity of stack is number of vertices.
    stack_push(s, v);

    while (!stack_empty(s)) {
        stack_pop(s, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                stack_push(s, u);
            }
        }
    }
}
```



Stack: 4


Top of stack will be the rightmost element.

Depth-First-Search (DFS)

```
void dfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Track where we've been.
    Stack *s = stack_create(graph_vertices(g)); // Capacity of stack is number of vertices.
    stack_push(s, v);

    while (!stack_empty(s)) {
        stack_pop(s, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                stack_push(s, u);
            }
        }
    }
}
```




Stack:

Top of stack will be the rightmost element.

Trees

- Trees are form of *acyclic* graph.
 - Acyclic means that if you follow the edges, there are *no loops* (cycles).
- You will often hear the term *DAG*, which stands for:
 - Directed
 - Acylic
 - Graph




Breadth-First-Search (BFS)

```
void bfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Capacity of queue is number of vertices.
    Queue *q = queue_create(graph_vertices(g)); // Capacity of queue is number of vertices.
    enqueue(q, v);

    while (!queue_empty(q)) {
        dequeue(q, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                enqueue(q, u);
            }
        }
    }
}
```




Queue: 0

Breadth-First-Search (BFS)

```
void bfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Capacity of queue is number of vertices.
    Queue *q = queue_create(graph_vertices(g)); // Capacity of queue is number of vertices.
    enqueue(q, v);

    while (!queue_empty(q)) {
        dequeue(q, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                enqueue(q, u);
            }
        }
    }
}
```




Queue: 1, 2, 3

Breadth-First-Search (BFS)

```
void bfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Capacity of queue is number of vertices.
    Queue *q = queue_create(graph_vertices(g)); // Capacity of queue is number of vertices.
    enqueue(q, v);

    while (!queue_empty(q)) {
        dequeue(q, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                enqueue(q, u);
            }
        }
    }
}
```




Queue: 2, 3, 4, 5

Breadth-First-Search (BFS)

```
void bfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Capacity of queue is number of vertices.
    Queue *q = queue_create(graph_vertices(g)); // Capacity of queue is number of vertices.
    enqueue(q, v);

    while (!queue_empty(q)) {
        dequeue(q, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                enqueue(q, u);
            }
        }
    }
}
```




Queue: 3, 4, 5, 6

Breadth-First-Search (BFS)

```
void bfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Capacity of queue is number of vertices.
    Queue *q = queue_create(graph_vertices(g)); // Capacity of queue is number of vertices.
    enqueue(q, v);

    while (!queue_empty(q)) {
        dequeue(q, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                enqueue(q, u);
            }
        }
    }
}
```




Queue: 4, 5, 6, 7

Breadth-First-Search (BFS)

```
void bfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Capacity of queue is number of vertices.
    Queue *q = queue_create(graph_vertices(g)); // Capacity of queue is number of vertices.
    enqueue(q, v);

    while (!queue_empty(q)) {
        dequeue(q, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                enqueue(q, u);
            }
        }
    }
}
```




Queue: 5, 6, 7

Breadth-First-Search (BFS)

```
void bfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Capacity of queue is number of vertices.
    Queue *q = queue_create(graph_vertices(g)); // Capacity of queue is number of vertices.
    enqueue(q, v);

    while (!queue_empty(q)) {
        dequeue(q, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                enqueue(q, u);
            }
        }
    }
}
```




Queue: 6, 7

Breadth-First-Search (BFS)

```
void bfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Capacity of queue is number of vertices.
    Queue *q = queue_create(graph_vertices(g)); // Capacity of queue is number of vertices.
    enqueue(q, v);

    while (!queue_empty(q)) {
        dequeue(q, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                enqueue(q, u);
            }
        }
    }
}
```




Queue: 7

Breadth-First-Search (BFS)

```
void bfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Capacity of queue is number of vertices.
    Queue *q = queue_create(graph_vertices(g)); // Capacity of queue is number of vertices.
    enqueue(q, v);

    while (!queue_empty(q)) {
        dequeue(q, &v);
        visited = set_insert(visited, v);

        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                enqueue(q, u);
            }
        }
    }
}
```




Queue: 7

Breadth-First-Search (BFS)

```
void bfs(Graph *g, uint32_t v) {
    Set visited = set_empty();           // Capacity of queue is number of vertices.
    Queue *q = queue_create(graph_vertices(g)); // Capacity of queue is number of vertices.
    enqueue(q, v);

    while (!queue_empty(q)) {
        dequeue(q, &v);
        visited = set_insert(visited, v);


        for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
            if (!set_member(visited, u) && graph_has_edge(g, v, u)) {
                enqueue(q, u);
            }
        }
    }
}
```



Queue:

Topological Sort

- Linearly orders vertices in a DAG such that, for any directed edge $\langle u, v \rangle$, u appears before v .
- Typically used to indicate ordering of dependencies.
 - Such as in a Makefile!
- Implemented using DFS:
 - Whenever a vertex is finished, it is prepended to a list.
 - Finished vertex: no more exploration possible from that vertex.
- Topological orderings are *not* unique.
 - Could be more than one valid topological ordering.



Modified DFS to Topologically Sort (Recursive)


```
void toposort(Graph *g, uint32_t v, Set *visited, LinkedList *order) {
    *visited = set_insert(*visited, v);

    for (uint32_t u = 0; u < graph_vertices(g); u += 1) {
        if (!set_member(*visited, u) && graph_has_edge(g, v, u)) {
            toposort(g, u, visited, order);
        }
    }

    ll_insert(order, v);
}

int main(void) {
    Graph *g = graph_create(8);
    graph_add_edge(g, 0, 1, 5); // <0, 1>, weight = 5
    graph_add_edge(g, 1, 2, 2); // <0, 1>, weight = 2
    graph_add_edge(g, 1, 3, 2); // <0, 1>, weight = 2
    graph_add_edge(g, 3, 2, 3); // <0, 1>, weight = 3
    graph_add_edge(g, 2, 4, 4); // <0, 1>, weight = 4

    Set visited = set_empty();
    LinkedList *order = ll_create();
    toposort(g, 0, &visited, order);
    return 0;
}
```



Order: 0 → 1 → 3 → 2 → 4

Kahn's Algorithm

Input: A directed acyclic graph $G = \langle V, E \rangle$.

Output: A list L containing a topological ordering of G .

KAHN(G)

```
1   $L \leftarrow []$  // List storing topological ordering.
2   $S \leftarrow \{\}$  // Set of vertices with no incoming edges.


---


3  for  $v \in V$ 
4      if  $v$  has no incoming edges
5           $S \leftarrow S \cup \{v\}$ 


---


6  while  $|S| > 0$ 
7      remove some  $v \in S$ 
8       $L \leftarrow L + [v]$ 
9      for each  $u \in V$  where  $e$  is  $\langle v, u \rangle$  and  $e \in E$ 
10          $E \leftarrow E - \{e\}$ 
11         if  $u$  has no incoming edges
12              $S \leftarrow S \cup \{u\}$ 


---


13 return  $L$ 
```

Single-Source Shortest Paths (SSSP)

- Assume some graph $G = \langle V, E \rangle$ and source vertex $s \in V$.
 - We want the shortest path from s to any $v \in V$.
- SSSP algorithms:
 - Bellman-Ford
 - Dijkstra's
 - We will focus on Dijkstra's algorithm.



Dijkstra's Algorithm

Input: A directed graph $G = \langle V, E \rangle$ and source vertex $s \in V$.


Output:

1. Distances from s for each vertex $v \in V$.
2. Predecessors for each $v \in V$ in shortest paths from s .

DIJKSTRA(G, s)

```
1   $S \leftarrow \{s\}$ 
2  for  $v \in G$ 
3       $D[v] \leftarrow \infty$  // Distance from source is infinite.
4       $P[v] \leftarrow -1$  // Predecessor of vertex is unknown.
5   $D[s] \leftarrow 0$  // Distance from  $s$  to  $s$ .
6  while  $S \neq V$ 
7       $w \leftarrow$  vertex in  $V - S$  such that  $D[w]$  is minimum
8       $S \leftarrow S \cup \{w\}$ 
9      for each  $v \in \text{ADJACENT}(G, w)$ 
10          $t \leftarrow D[w] + \text{WEIGHT}(G, w, v)$ 
11         if  $t < D[v]$  // If new distance from  $s$  to  $v$  is shorter.
12              $D[v] \leftarrow t$  // Update distance.
13              $P[v] \leftarrow w$  // Update predecessor.
14  return  $D, P$ 
```

$$D = [\overset{0}{\infty}, \overset{1}{\infty}, \overset{2}{\infty}, \overset{3}{\infty}, \overset{4}{\infty}]$$
$$P = [-1, -1, -1, -1, -1]$$
$$S = \{ \}$$



Dijkstra's Algorithm

Input: A directed graph $G = \langle V, E \rangle$ and source vertex $s \in V$.

Output:

1. Distances from s for each vertex $v \in V$.
2. Predecessors for each $v \in V$ in shortest paths from s .

DIJKSTRA(G, s)


```

1   $S \leftarrow \{s\}$ 
2  for  $v \in G$ 
3       $D[v] \leftarrow \infty$  // Distance from source is infinite.
4       $P[v] \leftarrow -1$  // Predecessor of vertex is unknown.
5   $D[s] \leftarrow 0$  // Distance from  $s$  to  $s$ .
6  while  $S \neq V$ 
7       $w \leftarrow$  vertex in  $V - S$  such that  $D[w]$  is minimum
8       $S \leftarrow S \cup \{w\}$ 
9      for each  $v \in \text{ADJACENT}(G, w)$ 
10          $t \leftarrow D[w] + \text{WEIGHT}(G, w, v)$ 
11         if  $t < D[v]$  // If new distance from  $s$  to  $v$  is shorter.
12              $D[v] \leftarrow t$  // Update distance.
13              $P[v] \leftarrow w$  // Update predecessor.
14  return  $D, P$ 
    
```

$$D = [\overset{0}{0}, \overset{1}{10}, \overset{2}{2}, \overset{3}{\infty}, \overset{4}{\infty}]$$

$$P = [-1, 0, 0, -1, -1]$$

$$S = \{0\}$$



$$\text{adjacent}(G, 0) = \{1, 2\}$$

vertex 1

$$t = 0 + 10$$

$$\text{is } 10 < \infty? \checkmark \text{ (update)}$$

vertex 2

$$t = 0 + 2$$

$$\text{is } 2 < \infty? \checkmark \text{ (update)}$$

Dijkstra's Algorithm

Input: A directed graph $G = \langle V, E \rangle$ and source vertex $s \in V$.

Output:

1. Distances from s for each vertex $v \in V$.
2. Predecessors for each $v \in V$ in shortest paths from s .

DIJKSTRA(G, s)


```

1   $S \leftarrow \{s\}$ 
2  for  $v \in G$ 
3       $D[v] \leftarrow \infty$  // Distance from source is infinite.
4       $P[v] \leftarrow -1$  // Predecessor of vertex is unknown.
5   $D[s] \leftarrow 0$  // Distance from  $s$  to  $s$ .
6  while  $S \neq V$ 
7       $w \leftarrow$  vertex in  $V - S$  such that  $D[w]$  is minimum
8       $S \leftarrow S \cup \{w\}$ 
9      for each  $v \in \text{ADJACENT}(G, w)$ 
10          $t \leftarrow D[w] + \text{WEIGHT}(G, w, v)$ 
11         if  $t < D[v]$  // If new distance from  $s$  to  $v$  is shorter.
12              $D[v] \leftarrow t$  // Update distance.
13              $P[v] \leftarrow w$  // Update predecessor.
14  return  $D, P$ 
    
```

$$D = [0, 6, 2, \infty, 7]$$

$$P = [-1, 2, 0, -1, 2]$$

$$S = \{0, 2\}$$



$$\text{adjacent}(G, 2) = \{1, 4\}$$

vertex 1

$$t = 2 + 4$$

$$\text{is } 6 < 10? \checkmark \text{ (update)}$$

vertex 4

$$t = 2 + 5$$

$$\text{is } 7 < \infty? \checkmark \text{ (update)}$$

Dijkstra's Algorithm

Input: A directed graph $G = \langle V, E \rangle$ and source vertex $s \in V$.

Output:

1. Distances from s for each vertex $v \in V$.
2. Predecessors for each $v \in V$ in shortest paths from s .

DIJKSTRA(G, s)

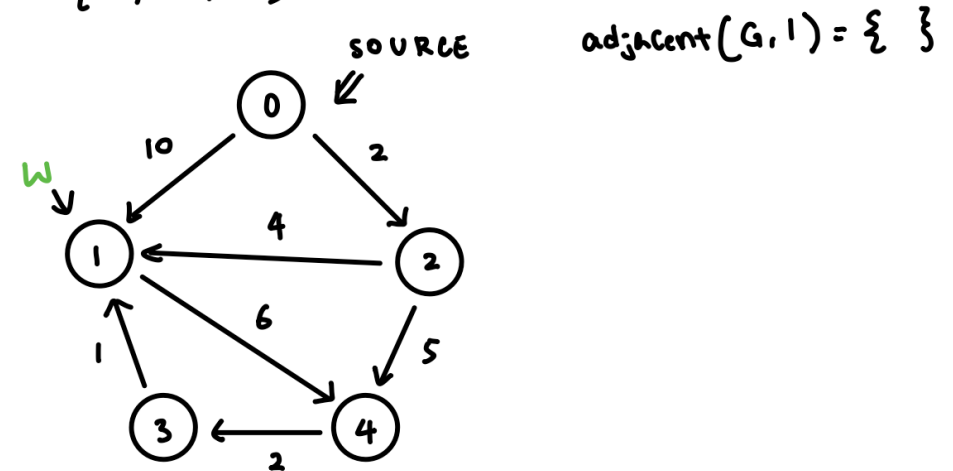
```

1   $S \leftarrow \{s\}$ 
2  for  $v \in G$ 
3       $D[v] \leftarrow \infty$  // Distance from source is infinite.
4       $P[v] \leftarrow -1$  // Predecessor of vertex is unknown.
5   $D[s] \leftarrow 0$  // Distance from  $s$  to  $s$ .
6  while  $S \neq V$ 
7       $w \leftarrow$  vertex in  $V - S$  such that  $D[w]$  is minimum
8       $S \leftarrow S \cup \{w\}$ 
9      for each  $v \in \text{ADJACENT}(G, w)$ 
10          $t \leftarrow D[w] + \text{WEIGHT}(G, w, v)$ 
11         if  $t < D[v]$  // If new distance from  $s$  to  $v$  is shorter.
12              $D[v] \leftarrow t$  // Update distance.
13              $P[v] \leftarrow w$  // Update predecessor.
14  return  $D, P$ 
    
```

$$D = [0, 6, 2, \infty, 7]$$

$$P = [-1, 2, 0, -1, 2]$$

$$S = \{0, 2, 1\}$$



Dijkstra's Algorithm

Input: A directed graph $G = \langle V, E \rangle$ and source vertex $s \in V$.

Output:

1. Distances from s for each vertex $v \in V$.
2. Predecessors for each $v \in V$ in shortest paths from s .

DIJKSTRA(G, s)


```

1   $S \leftarrow \{s\}$ 
2  for  $v \in G$ 
3       $D[v] \leftarrow \infty$  // Distance from source is infinite.
4       $P[v] \leftarrow -1$  // Predecessor of vertex is unknown.
5   $D[s] \leftarrow 0$  // Distance from  $s$  to  $s$ .
6  while  $S \neq V$ 
7       $w \leftarrow$  vertex in  $V - S$  such that  $D[w]$  is minimum
8       $S \leftarrow S \cup \{w\}$ 
9      for each  $v \in \text{ADJACENT}(G, w)$ 
10          $t \leftarrow D[w] + \text{WEIGHT}(G, w, v)$ 
11         if  $t < D[v]$  // If new distance from  $s$  to  $v$  is shorter.
12              $D[v] \leftarrow t$  // Update distance.
13              $P[v] \leftarrow w$  // Update predecessor.
14  return  $D, P$ 
    
```

$$D = [\overset{0}{0}, \overset{1}{6}, \overset{2}{2}, \overset{3}{9}, \overset{4}{7}]$$

$$P = [-1, 2, 0, 4, 2]$$

$$S = \{ 0, 2, 1, 4 \}$$



$\text{adjacent}(G, 4) = \{ 3 \}$

vertex 3

$t = 7 + 2$

is $9 < \infty$? \checkmark (update)

Dijkstra's Algorithm

Input: A directed graph $G = \langle V, E \rangle$ and source vertex $s \in V$.

Output:

1. Distances from s for each vertex $v \in V$.
2. Predecessors for each $v \in V$ in shortest paths from s .

DIJKSTRA(G, s)


```

1   $S \leftarrow \{s\}$ 
2  for  $v \in G$ 
3       $D[v] \leftarrow \infty$  // Distance from source is infinite.
4       $P[v] \leftarrow -1$  // Predecessor of vertex is unknown.
5   $D[s] \leftarrow 0$  // Distance from  $s$  to  $s$ .
6  while  $S \neq V$ 
7       $w \leftarrow$  vertex in  $V - S$  such that  $D[w]$  is minimum
8       $S \leftarrow S \cup \{w\}$ 
9      for each  $v \in \text{ADJACENT}(G, w)$ 
10          $t \leftarrow D[w] + \text{WEIGHT}(G, w, v)$ 
11         if  $t < D[v]$  // If new distance from  $s$  to  $v$  is shorter.
12              $D[v] \leftarrow t$  // Update distance.
13              $P[v] \leftarrow w$  // Update predecessor.
14  return  $D, P$ 
    
```

$$D = \begin{bmatrix} 0 & 6 & 2 & 9 & 7 \end{bmatrix}$$

$$P = \begin{bmatrix} -1 & 2 & 0 & 4 & 2 \end{bmatrix}$$

$$S = \{0, 2, 1, 4, 3\}$$



$\text{adjacent}(G, 3) = \{1, 4\}$

vertex 1


$t = 9 + 1$

is $10 < 6$? \times (no update)




Hamiltonian Path

- A path in an undirected or directed graph that visits each *vertex* exactly once.
 - Must start from an origin vertex and end up back at the origin.



Eulerian Path

- A path in an undirected or directed graph that visits each *edge* exactly once.
 - Must start from an origin vertex and end up back at the origin.



Summary

- Graphs pervade computer science.
 - Shortest path finding
 - Graph coloring
 - Network flow
 - Dependency ordering
 - ... and so much more!
- Come in undirected and directed forms.
- Used generally to indicate relationships between entities.
- Can be represented using either an adjacency matrix or using adjacency lists.