

# Unix File I/O

Dr. Kerry Veenstra

CSE 13S

May 19, 2023

# Using System Calls for File I/O

- See Section 2 of the man Pages
  - For example: `$ man 2 write`
- We won't be using Section 3 functions like `fopen()` or `printf()`.
  - Why? We want to control the buffering.

# Comparison: Buffered I/O vs. Unbuffered I/O

```
FILE *f = fopen(filename, mode);  
FILE *f = fopen("file.txt", "r");
```

vs.

```
int fd = open(filename, flags, permissions);  
int fd = open("file.txt", RD_RDONLY, 0);
```

```
int res = read(int fd, char *buf, int n) ;
```

- **fd** = file descriptor
  - 0 (**STDIN\_FILENO**)
  - result of **open()** or **creat()**
- **buf** = pointer to a buffer
- **n** = number of bytes to read
- **res** = number of bytes actually read (**res** < **n** is possible)
  - or 0 to indicate an end-of-file (EOF)
  - or -1 to indicate an error (see **errno**)

```
int res = write(int fd, char *buf, int n) ;
```

- **fd** = file descriptor
  - 1 (STDOUT\_FILENO) or 2 (STDERR\_FILENO)
  - result of **open()** or **creat()**
- **buf** = pointer to a buffer
- **n** = number of bytes to write
- **res** = number of bytes actually written (**res** < **n** is possible)
  - or **-1** to indicate an error (see **errno**)

```
int fd = open(char *name, int flags, int perms);
```

- **name** = file to open (file must exist, but see **creat()**)
- **flags** = O\_RDONLY, O\_WRONLY, or O\_RDWR
- **perms** = permissions code (usually 0)
- **fd** = file descriptor
  - or **-1** to indicate an error (see **errno**)

```
int fd = creat(char *name, int perms);
```

- **name** = file to open (overwrites if it exists)
  - Equivalent to `open(name, O_WRONLY|O_CREAT|O_TRUNC, perms)`
- **perms** = permissions code (also see `man chmod`)
- **fd** = file descriptor
  - or `-1` to indicate an error (see `errno`)

```
int res = close(int fd) ;
```

- **fd** = file descriptor
  - result of **open()** or **creat()**
- **res** = return code
  - **0** to indicate success
  - **-1** to indicate an error (see **errno**)



```
int res = unlink(char *name) ;
```

- **name** = name of file to delete
- **res** = return code
  - 0 to indicate success
  - -1 to indicate an error (see errno)

```
int res = lseek(int fd, long offset, int origin);
```

- **fd** = file descriptor
  - result of **open()** or **creat()**
- **offset** = offset in bytes relative to the origin
- **origin**
  - **SEEK\_SET** = 0 = beginning of file
  - **SEEK\_CUR** = 1 = current position
  - **SEEK\_END** = 2 = end of file
- **res** = return code
  - 0 to indicate success
  - -1 to indicate an error (see **errno**)

# Buffering Example

```
#include <stdio.h>
#include <unistd.h>

/*
 * Demonstration that combining buffered I/O like fprintf()
 * and unbuffered I/O like write() can lead to out-of-order
 * data in the output file.
 */
int main(void) {
    FILE *f = fopen("out.txt", "w");
    int fd = fileno(f);
    fprintf(f, "1\n");
    write(fd, "2\n", 2);
    fprintf(f, "3\n");
    write(fd, "4\n", 2);
    fclose(f);
    return 0;
}
```

# Buffering Example

```
#include <stdio.h>
#include <unistd.h>

/*
 * Demonstration that combining buffered I/O like fprintf()
 * and unbuffered I/O like write() can lead to out-of-order
 * data in the output file.
 */
int main(void) {
    FILE *f = fopen("out.txt", "w");
    int fd = fileno(f);
    fprintf(f, "1\n");
    write(fd, "2\n", 2);
    fprintf(f, "3\n");
    write(fd, "4\n", 2);
    fclose(f);
    return 0;
}
```

# Buffering Example

```
#include <stdio.h>
#include <unistd.h>

/*
 * Demonstration that combining buffered I/O like fprintf()
 * and unbuffered I/O like write() can lead to out-of-order
 * data in the output file.
 */
int main(void) {
    FILE *f = fopen("out.txt", "w");
    int fd = fileno(f);
    fprintf(f, "1\n");
    write(fd, "2\n", 2);
    fprintf(f, "3\n");
    write(fd, "4\n", 2);
    fclose(f);
    return 0;
}
```

# Buffering Example

```
[veenstra@arm128:~/130-CSE$ clang -o buffering bufferi
ng.c
[veenstra@arm128:~/130-CSE$ ./buffering
[veenstra@arm128:~/130-CSE$ cat out.txt
2
4
1
3
veenstra@arm128:~/130-CSE$
```

# Buffering Example

```
veenstra@arm128:~/130-CSE$ clang -o buffering buffering.c
veenstra@arm128:~/130-CSE$ ./buffering
veenstra@arm128:~/130-CSE$ cat out.txt
2
4
1
3
veenstra@arm128:~/130-CSE$
```